

Reminders and Topics

- **Search project due Friday**
 - Last project, Yay!
- **SRTI Course Survey - online**
 - Survey period: 12/3 7am - 12/20 11pm
- This lecture:
 - **Hashing and Hash Tables**
- Next lecture:
 - **Final review**

So Far We've Learned

data structure	add	search	remove
unsorted array	$O(1)$	$O(N)$	$O(N)$
unsorted linked	$O(1)$	$O(N)$	$O(N)$
sorted array	$O(N)$	$O(\log N)$	$O(N)$
BST (balanced / worst)	$O(\log N) / O(N)$	$O(\log N) / O(N)$	$O(\log N) / O(N)$
heap	$O(\log N)$	$O(N)$	$O(\log N)$

Is this the best? Can we do better??

Hash Table

- **Surprisingly Fast**
 - On average, $O(1)$ for all of add/search/delete!
- **Surprisingly Simple**
 - Easy to implement
- **No way, it can't perfect. What's the catch?**
 - Need to have a good idea about the number of elements. Difficult to re-size dynamically.
 - Performance degrades when it's close to full.
 - Can't easily visit data in sorted order (such as getting the max).

A First Idea

- Here is a trivial way to achieve $O(1)$ cost: say we want to store and search student records based on student ID, and we assume each ID is unique.
- Since each ID is 8-digit long, let's prepare an array / table with a capacity of 100,000,000, so there is one entry reserved for each possible ID (0 to 99,999,999).
- Add / Search / Remove all cost just $O(1)$!
- What's the problem with this approach?

Not all 8-digit integers are valid student IDs!

Many entries are empty. This is a huge waste of storage!


Hashing



- UMass has about 22,000 undergrads (way less than the number of 8-digit integers). Ideally we want a table that's not too much larger than that.
- How do we know where each ID should be stored at?
- Hashing comes to rescue:
 - What does the word **hash** mean? Hash browns?
 - Hashing maps a key value (which can span a wide range) to an index (which has a much smaller range).
 - Sparsity is characteristic of such data
 - credit card numbers, dictionary words

Hashing

- How do we map the huge range of possible key values to a much smaller range, so that they can be stored in a reasonably sized array?
 - Example: map an 8-digit key to 50,000 size array. There can be many different ways, what's the simplest?
- Use **modulo (%)**: **$\text{index} = \text{key} \% \text{array_size}$**
- This is called a **hash function**.
- **array_size** must be at least the number of elements (e.g. the number of students), but often a few times bigger



Collisions

- One obvious problem is that multiple keys can map to the same index (below array_size is 50000):
 - $23245467 \% 50000 = 45467$
 - $43345467 \% 50000 = 45467$
 - In fact, $(45467 + 50000 \times k) \% 50000 = 45467$ for any positive integer k !
- This is called **collision**. It can be reduced by using a better hash function (e.g. array size should always be a prime number). But it cannot be avoided.

Collision Handling

- There are empty slots in the hash table to store collided elements, because we know the array size is at least the number of elements.
- We need a systematic way to search for such empty slots when collision happens. There are two generally approaches:
 - **Open addressing**
 - **Separate chaining**

Open Addressing

- **Linear Probing**

When **inserting** a new element, if a collision is encountered, check the succeeding slots one by one, until an empty slot is found.


$\text{index}, \text{index}+1, \text{index}+2, \dots$ (the array is circular, so back to 0 when reaching the end of the array)

- Guaranteed to find one empty slot eventually.
- Demo.
- Probe length (i.e. step size) doesn't necessarily have to be 1.


Linear Probing

For now, assume we don't remove any elements yet.

When **searching** for an element, probe linearly until either the element is found, **or an empty slot is encountered** (i.e. element does not exist). Why?



a) Successful search for 472



b) Unsuccessful search for 893

Linear Probing

- What about **deletion**?
 - Unlike before, you can't delete an element by simply setting it to null. This will affect how you search. Why?
- How do we solve this problem?
 - Solution: **lazy deletion**. We don't set the deleted element to null, instead, we set a special flag to indicate it's gone. Insertion can overwrite a flagged slot; but search must continue across it.

Lazy Delete Example

(index = key % 10)

Insert 11, 21, then 32

0	1	2	3
null	null	null	null

0	1	2	3
null	11	21	32

Delete 21 (it gets **flagged**)
Search can still find 32

0	1	2	3
null	11	21	32

Insert 31

0	1	2	3
null	11	31	32

Generic Hash Table

```
class DataEntry<T>{
    public int key; // As you will see later key can also be generic...
    public T value;
    public boolean flag;
    private DataEntry(int insertKey, T insertValue){
        key = insertKey;
        value = insertValue;
        flag = false;
    }
}

class HashTable<T> {
    protected DataEntry<T>[] table;
    protected int numElements=0;
    final int DEFCAPACITY = 100;
    public HashTable(int capacity) {
        table = (DataEntry<T>[]) new Object[DEFCAPACITY];
    }
}
```

Linear Probing

```
// assume elements are stored in table  
// the capacity of the table is table.Length;  
public T search (int key){  
    int index = key % table.Length; // hash func  
    int start = index; // keep a copy of index  
    while(table[index] != null) {  
        if(table[index].key == key && !table[index].flag)  
            return table[index].value;  
        index = (index+1) % table.Length; // what's this?  
        if(index==start) return null; // if back to start  
    }  
    return null;  
}
```

Wrap around to the beginning if you've reached the end of the table. This is just like the circular array queue.

Linear Probing

```
public T delete (int key){  
    int index = key % table.length; // hash func  
    while(table[index] != null)  
        if(table[index].key == key && !table[index].flag) {  
            T value = table[index].value;  
            table[index].flag = true;  
            return value;  
        }  
        index = (index+1) % table.length;  
    }  
    return null;  
}
```

Caveat: What happens when the item doesn't exist?

Linear Probing

```
public void insert(int key, T value){  
    int index = key % table.length; // hash func  
    while(table[index] != null &&  
        !table[index].flag)  
        index = (index+1) % table.length;  
  
    // new item always has a cleared flag  
    table[index] = new DataItem(key, value);  
}
```

Caveat: what happens when the table is full?

Clicker #1

The following hashtable has a **capacity of 10**. The index of each slot is shown below for convenience. Using **linear probing** (probe length=1), where would a new element **59** be inserted into?

60			53	13	25		27		48
[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]	[8]	[9]

A. [1]

B. [2]

C. [6]

D. [8]

E. [9]

Linear Probing

```
// assume elements are stored in table  
// the capacity of the table is table.length;  
public T search (int key){  
    int index = key % table.length; // hash func  
    int start = index; // keep a copy of index  
    while(table[index] != null) {  
        if(table[index].key == key && !table[index].flag)  
            return table[index].value;  
        index = (index+probe_length) % table.length;  
        if(index==start) return null; // if back to start  
    }  
    return null;  
}
```

Probe length doesn't have to be 1. It can be any positive integer of your choice.

Clicker #2

The following hashtable has a **capacity of 10**. The index of each slot is shown below for convenience. Using **linear probing** (**probe length=4**), where would a new element **59** be inserted into?

60			53	13	25		27		48
[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]	[8]	[9]

A. [1]

B. [2]

C. [6]

D. [8]


E. [9]

Linear Probing

- **Clustering problem**

Search can take a long time if a cluster is formed.

What happens when the array is close to full?




Quadratic Probing

When collision happens, check the succeeding slots in **quadratic** steps / probe lengths:


index , $\text{index}+1$, $\text{index}+4$, $\text{index}+9$, $\text{index}+16$,
 $\text{index}+25$...

- Using increasingly larger steps helps reduce the possibility of forming clusters.

Quadratic probing



a) Successful search for 420



b) Unsuccessful search for 481

Double Hashing

- The problem with linear and quadratic probing is that once keys collide, they all follow exactly the same probing path, completely predictable.
- We need a way to generate probe lengths that can vary and are not pre-determined.
- The solution is to pick a **probe length that varies depending on the key value**. This can be achieved using a second hashing function, thus the name 'double hashing'.

Double Hashing

- A good choice for secondary hashing:


$\text{probe_length} = \text{constant} - (\text{key} \% \text{constant});$

- What's the range of possible values?

- For example:


$\text{probe_length} = 5 - (\text{key} \% 5);$

- Thus when different keys hash to the same index, they will most likely generate different probe lengths
- Note that this secondary hashing function never generates a probe len of 0. Why is this important?



Key=888

Probe length = $[5 - 888 \% 5] = 2$



Key=710

Probe length = $[5 - 710 \% 5] = 5$

Clicker #3

The following hashtable has a **capacity of 10**. The index of each slot is shown below for convenience. Using **double hashing** where the probe length is calculated as $[7 - (\text{key} \% 7)]$, where would a new element **80** be inserted into?

60			53	13	25		27		48
[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]	[8]	[9]

A. [0]

B. [1]

C. [2]

D. [6]

E. [8]

Load Factor and Rehashing

- **Load Factor**: number of elements divided by the hash table size / capacity.
 - Example: if a table has a capacity of 73, currently storing 40 elements. The load factor is:
 $40 / 73 = 56\%$
- When load factor becomes too large (close to 1, i.e. table getting full), it's necessary to increase the hash table capacity. This is called **re-hashing**.

Clicker #4

```
void rehashing() {  
    T[] nheap = (T[]) new Object[heap.length * 2];  
    for (int i=0; i < heap.length; i++) {  
        Nheap[i] = heap[i];  
    }  
    heap = nheap;  
}
```

What's wrong of this rehashing method?

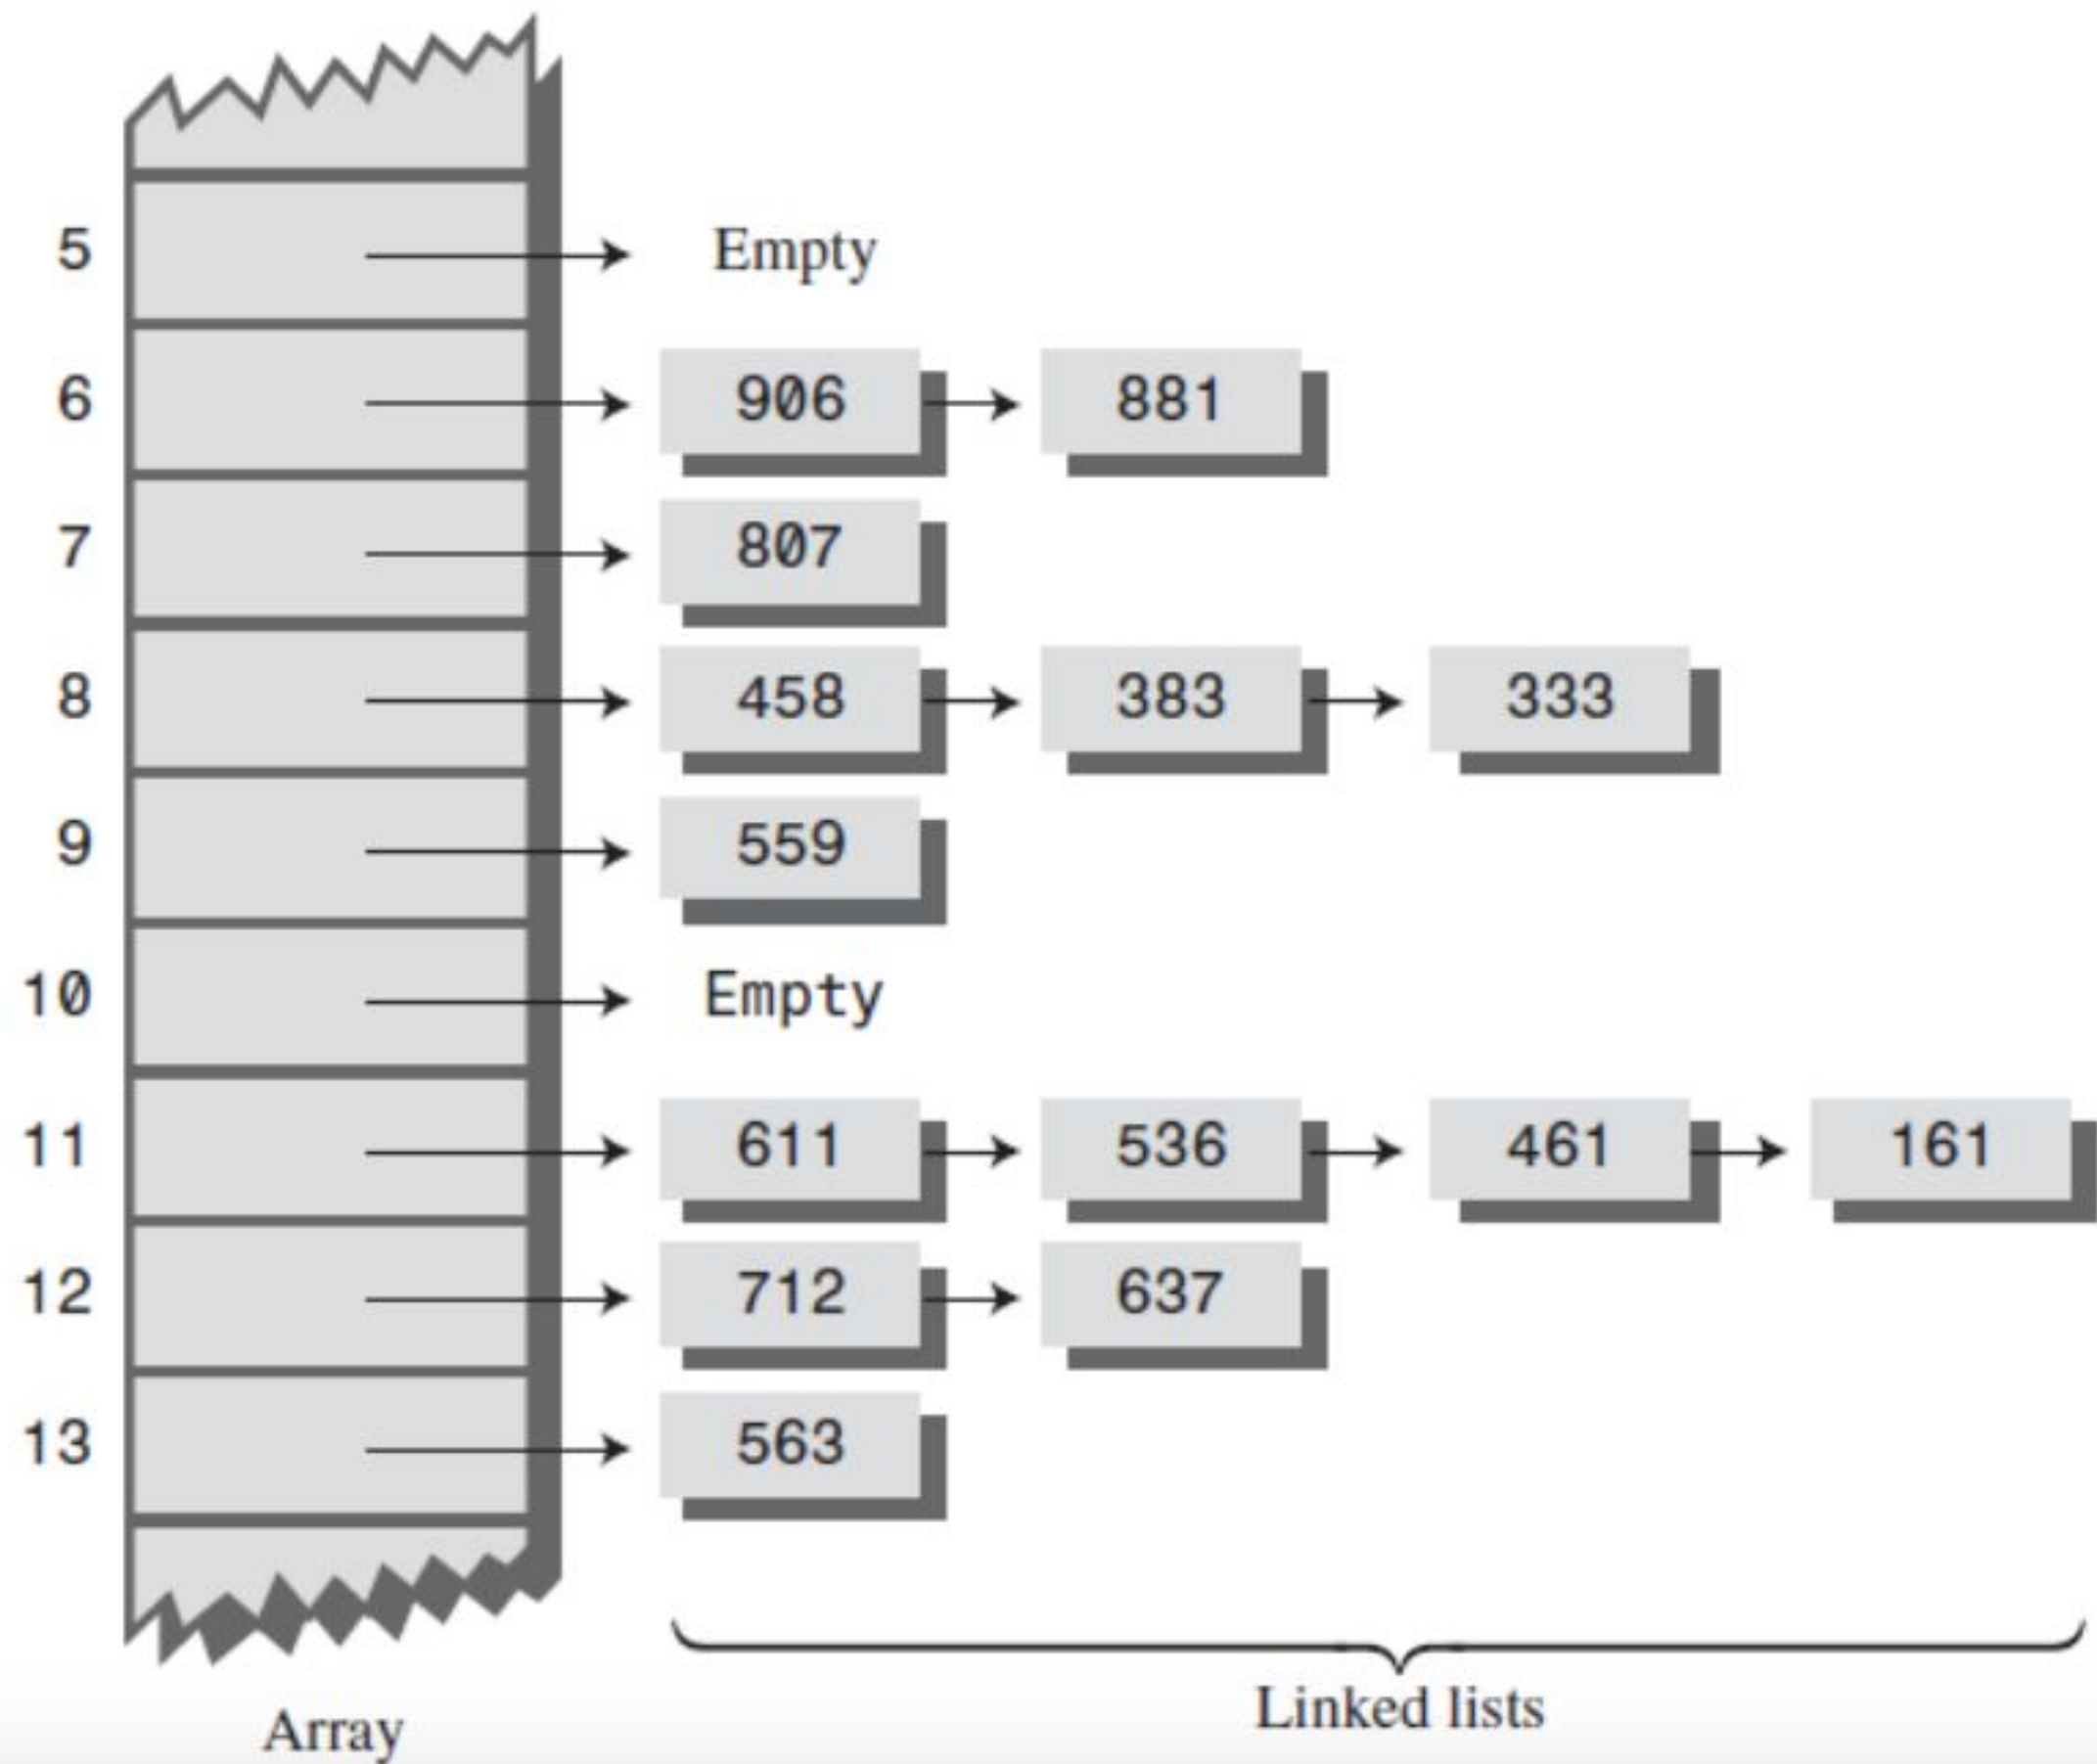
- 1) It works just fine.
- 2) It doesn't enlarge the heap space.
- 3) It doesn't copy the elements to the new heap.
- 4) It doesn't update hash value of the elements for new heap

Load Factor and Rehashing

- **Load Factor**: number of elements divided by the hash table size / capacity.
 - Example: if a table has a capacity of 73, currently storing 40 elements. The load factor is:
 $40 / 73 = 56\%$
- When load factor becomes too large (close to 1, i.e. table getting full), it's necessary to increase the hash table capacity. This is called **re-hashing**.
- Since the table size has changed, you need to rehash every element to its new location in the new hash table. You can't simply copy elements over. Why?

Separate Chaining

- Open Addressing (e.g. linear, quadratic probing and double hashing) resolves collision by looking for an empty slot in the hash table.
- Another idea is to create a linked list at each slot to allow for multiple elements. This is called **Separate Chaining**.
- In separate chaining, deletion is easier because all collided elements go into a linked list, so deletion simply deletes the element from that linked list. No 'flag' business any more.



Java's Hashtable class

- Java provides a class `Hashtable<K, V>` where K is key's type, V is the value's type
- It implements the `Map<K, V>` interface, and can re-hash dynamically based on a pre-defined load factor (e.g. 0.75)

- For example:

```
Hashtable<String, Float> table =  
    new Hashtable<String, Float>();
```

```
table.put("John Smith", 6.0f);
```

```
table.put("Eric May", 5.8f);
```

```
table.put("Rose Ann", 5.9f);
```

```
System.out.println(table.get("Rose Ann"));
```

Questions?